

LINEAR REGRESSION III

11.19.2018

HOMEWORK 4

- * due a week from today!
- * there will be an extra office hour
TOMORROW (Tuesday, 11/20/2018) from
10-11am

RECAP

- * `np.linalg.lstsq` – numpy function that does least squares regression
- * R-squared is a measure of how good a regression model is

RECAP

- * in-set vs. out-of-set evaluation of a regression model
- * in-set biased upwards (to 1.0 in even not-so-extreme cases!)
- * out-of-set unbiased (maybe?) or biased downwards, depending on assumptions

RECAP

- * regression stability
- * similar (or linearly dependent) regressors can cause instability in the weight estimates
- * can be assessed by looking at the **singular values** output by **lstsq**
- * tiny singular values correspond to unstable directions in regression

STABILITY

- * how do we correct unstable regression?
- * remember the issue is that many different values for the weights can give us ~the same answer

STABILITY

- * we correct unstable regression by making *assumptions* about what the weights should look like
- * this is called **regularization**
 - * *imo this is the most important concept in all of machine learning*

STABILITY

- * the most common assumption is that the weights should be *small*
- * what does “small” mean? and how can we enforce “smallness”?

REGULARIZATION

- * one way to get regularization is to modify the **error function** that we minimize to get the weights
- * we want small weights, so we can make large weights look like an error!

REGULARIZATION

- * recall the “least squares” error function:

$$Err(\beta) = \sum_{t=1}^T (y_t - x_t \beta)^2$$

REGULARIZATION

$$\text{Err}(\beta) = \sum_{t=1}^T (y_t - x_t \beta)^2 + \lambda \sum_{i=1}^P \beta_i^2$$

* we can modify this to add a *penalty* for large weights

$$\text{Err}(\beta) = \sum_{t=1}^T (y_t - x_t \beta)^2 + \lambda \sum_{i=1}^P \beta_i^2$$

REGULARIZATION

* now the error is the sum of a **loss term** and a **penalty term**

$$Err(\beta) = \sum_{t=1}^T (y_t - x_t \beta)^2 + \lambda \sum_{i=1}^P \beta_i^2$$

REGULARIZATION

- * it also introduces an extra parameter, λ , which is the *regularization coefficient*, or, in this case, *ridge coefficient*

$$Err(\beta) = \sum_{t=1}^T (y_t - x_t \beta)^2 + \lambda \sum_{i=1}^P \beta_i^2$$

REGULARIZATION

- * what effect does adding this type of regularization have on a regression model?

1D EXAMPLE

subject



stimulus



- * y = output of a neuron that you are measuring
- * x = how many times per second the screen flashes

1D EXAMPLE


subject




stimulus




y




x



1D EXAMPLE





1D EXAMPLE



1D EXAMPLE


$$\text{Err}(\beta) = \sum_{t=1}^T (y_t - x_t \beta)^2$$

$$\text{Err}(\beta) = \sum_{t=1}^T (y_t - x_t \beta)^2$$




1D EXAMPLE

$$\text{Err}(\beta) = \sum_{t=1}^T (y_t - x_t \beta)^2$$



$$\text{Err}(\beta) = \sum_{t=1}^T (y_t - x_t \beta)^2$$




$$Err(\beta) = \sum_{t=1}^T (y_t - x_t \beta)^2 + \lambda \beta^2$$

1D EXAMPLE

Regularization:

$$Err(\beta) = \sum_{t=1}^T (y_t - x_t \beta)^2 + \lambda \beta^2$$




$$\text{Err}(\beta) = \sum_{t=1}^T (y_t - x_t \beta)^2 + \lambda \beta^2$$

1D EXAMPLE

Regularization:

$$\text{Err}(\beta) = \sum_{t=1}^T (y_t - x_t \beta)^2 + \lambda \beta^2$$



REGULARIZATION

- * what effect does adding this type of regularization have on a regression model?
- * it changes the shape of the error function to be more *circular* (and thus more stable)
- * (example)

REGULARIZATION

- * what effect does this have on the weights?
- * among the many \sim equivalent sets of weights, it preferentially chooses those that are *small*
- * (example)

RIDGE REGRESSION

- * this type of regularization (penalizing the sum of squared weights) is called **ridge regression**
- * and because the ridge error function (loss + penalty) is parabolic, it has an analytic solution!

RIDGE REGRESSION

- * a nice implementation is in scikit-learn (a package that we'll talk more about next week) in `sklearn.linear_model.Ridge`

RIDGE REGRESSION

- * but when doing ridge regression you have a new issue: how do you choose the ridge parameter, λ ?

RIDGE REGRESSION

- * if you train and test your regression model on the same piece of data, $\lambda=0$ is always going to be the best
- * *~bogus~*

RIDGE REGRESSION

- * if you train and test on different datasets (as discussed friday) it's better
- * but using your test data multiple times (to choose a parameter!) creates an issue of bias (aka **overfitting**)

RIDGE REGRESSION

- * the correct solution is **cross-validation**:
- * break your dataset into training and test ($X \rightarrow [X_{\text{trn}}, X_{\text{test}}]$)
- * further break up your training set ($X_{\text{trn}} \rightarrow [X_{\text{fit}}, X_{\text{val}}]$)
- * fit weights using X_{fit} , choose λ based on performance on X_{val} , then finally test on X_{test}

END